



Coruscant: Co-Designing GPU Kernel and Sparse Tensor Core to Advocate Unstructured Sparsity in Efficient LLM Inference

Donghyeon Joo¹, Helya Hosseini¹, Ramyad Hadidi², Bahar Asgari¹


¹University of Maryland, College Park

²d-Matrix

LLMs are huge and memory-hungry

Various model compression techniques are used

Distillation




Small model learns
from larger model



DeepSeek-R1
[2025]

Quantization




Reduce model bit-width



gpt-oss
[2025]

Pruning




Reduce neural connections

?

LLMs are huge and memory-hungry

Various model compression techniques are used

Distillation




Small model learns
from larger model



DeepSeek-R1
[2025]




Quantization



Reduce model bit-width




gpt-oss
[2025]




Tradeoff in Model Accuracy and Efficiency


Structured Sparsity



Semi-Structured Sparsity




Unstructured Sparsity




Tradeoff in Model Accuracy and Efficiency

Structured Sparsity



Smaller dense matrix


Semi-Structured Sparsity



position
metadata

Constant number of
non-zeros

Unstructured Sparsity



Compute as dense

Efficient Hardware Mapping

Accuracy Retention

Tradeoff in Model Accuracy and Efficiency

Methods [†] (pattern)	LLM	Perplexity Diff. [‡]
SparseGPT (unstructured)	OPT-175B	-0.14
SparseGPT (2:4)	OPT-175B	0.39
Wanda (unstructured)	Llama-65B	1.01
Wanda (2:4)	Llama-65B	2.69


Perplexity on WikiText-2.

Perplexity difference from dense model, lower the better


	25% (3:4)	37.5% (3:8)	50% (2:4)	62.5 % (5:8)	75% (1:4)
Semi-Structured	84.22	75.70	6.09	0.06	0.00
Unstructured	88.40	84.43	42.81	2.21	0.43

TriviaQA accuracy on Llama-2 7B
Pruned to each sparsity with Wanda.


Effective Sparsity Range in LLM Pruning



Efficient Unstructured Sparsity in Hardware




- We target a moderately sparse region.
- Sparse formats with numerical non-zero positions fail to compress our target.




Flash-LLM sparse format at 50% sparsity

Efficient Unstructured Sparsity in Hardware



- We target a moderately sparse region.
- Sparse formats with numerical non-zero positions fail to compress our target.
- Bitmap-based representation is the only plausible solution for compression.


Coruscant




Coruscant Sparse Format

- Bitmap-based non-zero position representation
- GPU Kernel and Architecture considerations:
 - Each warp thread assigned to two tiles
 - Column-wise tiling to avoid shared memory bank conflict
- Maximal Compression Benefits:
 - Accelerated memory-bound SpMM
 - Reduced global memory consumption


LLM Weight (float16)				
T1	T2	T3	T4	
30 NNZ	5 NNZ	4 NNZ	8 NNZ	
...	
...	
...	



Coruscant




Coruscant GPU Kernel - Target Operation



- Inference is dominated by decode phase.
- Decode phase is dominated by weight projection SpMM.
- Due to small batch dimension, weight projection SpMM is memory bound.

Latency Comparison of Prefill and Decode

Coruscant GPU Kernel - Target Operation



Coruscant GPU Kernel - Load-Compute Pipeline

- Objective: Keep the data transfer from GMEM to GPU processor compressed.
- Phase 1: Matrix tiles are transferred to processor registers as-compressed, then decompressed to shared memory.




Figure credit to
Flash-LLM (Xia et al. VLDB 2023)

Coruscant GPU Kernel - Load-Compute Pipeline

- Phase 2: Decompressed matrix tiles are computed in Tensor Core.
- Each iteration is an overlapped execution of load and compute.




Figure credit to
Flash-LLM (Xia et al. VLDB 2023)


Coruscant GPU Kernel - Decompression Logic

- Decompression loop using count-leading-zero instructions to find set bits quickly.
- Column-wise tiling of Coruscant sparse format prevents shared memory bank conflict.

Algorithm 1 Bitmap decompression algorithm of a tile


```
1: for  $i \leftarrow 0$  to 63 do
2:   if  $i = nnz\_tile$  then
3:     break
4:   end if
5:    $pos1 \leftarrow \text{clz}(bmp)$ 
6:    $mask \leftarrow (1 \ll (63 - pos1))$ 
7:    $bmp \leftarrow bmp \& \sim mask$ 
8:    $output\_idx \leftarrow tile\_idx + (pos1 \ll 6)$ 
9:    $SmemPTR[output\_idx] \leftarrow Reg\_NZs[i]$ 
10:  end for
```

Variable loop iteration (i.e. number of NZs) cause register spilling.



Comparing Tiling Schemes on SMEM Write

Coruscant



Coruscant Sparse Tensor Core - Motivation

- Software decompression takes up to 36% of kernel execution time.
- In each load-compute pipeline iteration, decompression adds latency.






Figure credit to
Flash-LLM (Xia et al. VLDB 2023)


Coruscant Sparse Tensor Core

- Lightweight Bitmap Decoder to perform hardware decompression.
 - Decodes the thread-local bitmap to stream non-zeros to compute lanes.





Warp-Level MMA Formulation




Per-thread Bitmap Decoder

Evaluation Methodology


- **System:** NVIDIA RTX 6000 ADA GPU
- **STC Simulation Methodology:** Coruscant STC modeled as a kernel without decompression, but equal number of MMA instructions.
- **Key Metrics:**
 - **Efficiency:** SpMM/end-to-end latency, GMEM consumption, tensor core utilization
 - **LLM Quality:** Perplexity on Wiki-Text, TriviaQA accuracy from Longbench
- **Workloads:**
 - SpMM evaluation: Weight projection matrices of OPT 13B and 66B
 - End-to-end inference: Llama-2 7B and 13B

Evaluation: Coruscant Kernel



- Coruscant Kernel outperforms dense and sparse SOTAs across target sparsity.
- Outperforms dense cuBLAS with 1.09x (at 30%) to 2.00x (at 70%) speedup
- Outperforms sparse FlashLLM with 1.05x (at 70%) to 1.48x (at 30%) speedup


Evaluation: Coruscant Sparse Tensor Core



- Hardware decompression achieves average 1.3x speedup over Coruscant Kernel.
- Removing decompression latency improves Tensor Core utilization.
- Reducing shared memory read/write improves shared memory stalls.



Coruscant in Roofline Analysis

- Speedup is primarily attributed to bitmap compression reducing data transfer.




Coruscant in Roofline Analysis

- Speedup is primarily attributed to bitmap compression reducing data transfer.
- In roofline, hardware decompression represents an upright movement.



Evaluation: End-to-End Inference

- Llama-2-7B: Up to 26% increase in token generation throughput.
- Llama-2-13B: Up to 40% increase in token generation throughput.
- Reduces global memory footprint to support larger batch size.



Additional Evaluations

- Please refer to our paper for evaluation and analysis on ...
 - Latency, power, and area comparison with previous STC designs
 - Sparsity exploitation characterization: compute-skipping versus compression
 - Comparison between Coruscant and 2:4 semi-structured sparsity kernel
 - Comparison between Coruscant and N:M semi-structured accelerators
 - Performance on different workloads (Summarization, Conversational, Reasoning)
 - Impact on weight transfer over PCIe.
 - and more!

Conclusion

- Coruscant is a **software-hardware co-design** that keeps unstructured sparse weight matrices compressed from global memory to tensor core execution.
- **Contributions:**
 - ✓ A **bitmap-based sparse format** for compressing unstructured sparsity
 - ✓ A **GPU kernel** to perform *load-as-compressed, compute-as-dense*
 - ✓ A **sparse tensor core** to eliminate decompression latency
- Demonstrates that **unstructured sparsity can be practical and efficient** on modern GPUs with minimal hardware cost.

Thank You

We are happy to take
Questions and Comments

dhjoo98@umd.edu



This work is supported by
National Science Foundation (NSF).



Backup Slides

Coruscant Sparse Tensor Core - Bitmap Decoder

- Warp-level MMA formulation is altered to stream compressed non-zero from register to TC.
- Consecutive 4-bits are decoded every cycle to stream a decompressed 4-element vector.



Analysis: Compute Skipping vs Reducing Data Movement

- Two ways to leverage sparsity:
 - Compute Skipping accelerates computation.
 - Sparse matrix compression reduces memory footprint.
 - Previous STC designs for unstructured sparsity prioritized Compute Skipping.
 - Architecture for semi-structured sparsity shows strength in both.
 - Coruscant prioritizes maximal compression, targeting memory-bound SpMM.





Figure credit to
Flash-LLM (Xia et al. VLDB 2023)

Analysis: In-depth Comparison with Previous STCs

- Previous STC designs (DSTC, RM-STC) for unstructured sparsity seek to skip computation.
- Smaller tile size leads to 6%/12% non-zero pointer size, deteriorating compression ratio.
- This directly impacts performance on memory-bound SpMM.
- Compute-skipping increase the complexity of accumulation logic, increasing HW overhead.




	Added Area	Percentage to GPU	Power
Volta (V100)	0.51 mm ²	0.006%	26.24 mW
Ampere (A100)	0.68 mm ²	0.008%	35.43 mW
Hopper (H100)	1.44 mm ²	0.018%	74.79 mW

Analysis: Comparison with Semi-Structured Sparsity

- Efficiency - model accuracy trade-off applies to various N:M patterns for different sparsity.
- Semi-structured sparsity can compress better, but model accuracy degrades significantly.
- Coruscant frees future pruning works from sparsity pattern constraints



Table 4: TriviaQA score of semi-structured and unstructured sparsity at various sparsity levels.

	25% (3:4)	37.5% (3:8)	50% (2:4)	62.5 % (5:8)	75% (1:4)
Semi-Structured	84.22	75.70	6.09	0.06	0.00
Unstructured	88.40	84.43	42.81	2.21	0.43



Analysis: In-depth Comparison with 2:4 cuSPARSELt

- cuSPARSELt leverages the hardware decompression support from NVIDIA STC.
- Compared to Coruscant Kernel-only, cuSPARSELT dominates performance and TC utilization.
- Coruscant STC shows comparable performance to 2:4 semi-structured sparsity, while being flexible to arbitrary sparsity.



Evaluation: Generality and Workload Types

- Due to low occupancy, Coruscant underperforms for compute-bound prefill.
- Across representative input-output token ratio, slow prefill is amortized.

